

## summary

Yifu Yuan

2020/3/12

The following is the result of the model I trained last time. It has high loss and variance, with a low accuracy of 69%.




Figure 1: Last Week acc 69%

This week I tried two versions of models. The following is the one with four layers of CNN, with accuracy 78%. I increased the number of tensors in some layers and had different dropout value in each layer. See figure 2.

Figure 3 is the result of the other version. It has six layers of CNN, with dropout, L2 regularization, with accuracy 84%. However, I did not succeed in reducing the loss and variance of this version. I am not sure if it is the problem of learning rate, because I set learning rate to decay over time.




Figure 2: Version 1 accuracy 78%

The code of learning rate:

```
global_step = tf.Variable(0, trainable=False)

learning_rate = tf.compat.v1.train.exponential_decay(0.001, global_step, 1000, 0.9)

//l2 is the list of l2-regularization of each layer

cost = tf.reduce_mean(tf.losses.softmax_cross_entropy(logits=logits, onehot_labels=y)) + l2[0] + l2[1] +
l2[2] + l2[3] + l2[4] + l2[5]

optimizer = tf.compat.v1.train.AdamOptimizer(learning_rate=learning_rate).minimize(cost,global_step)
```

The code of CNN is here:

```
def conv_net(x,keep_prob):

    #L2 regularization

    l2 = []

    conv1_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 3, 32], mean=0, stddev=0.08))
    conv2_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 32, 32], mean=0, stddev=0.08))
    conv3_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 32, 64], mean=0, stddev=0.08))
    conv4_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 64, 64], mean=0, stddev=0.08))
    conv5_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 64, 128], mean=0, stddev=0.08))
    conv6_filter = tf.Variable(tf.truncated_normal(shape=[3, 3, 128, 128], mean=0, stddev=0.08))

    conv1 = tf.nn.conv2d(x, conv1_filter, strides=[1,1,1,1], padding='SAME')
    conv1_l2 = 1e-4*tf.nn.l2_loss(conv1)
    l2.append(conv1_l2)
    conv1 = tf.nn.relu(conv1)
    conv1_bn = tf.layers.batch_normalization(conv1)

    conv2 = tf.nn.conv2d(conv1_bn, conv2_filter, strides=[1,1,1,1], padding='SAME')
    conv2_l2 = 1e-4*tf.nn.l2_loss(conv2)
    l2.append(conv2_l2)
    conv2 = tf.nn.relu(conv2)
    conv2_pool = tf.nn.max_pool(conv2, ksize=[1,2,2,1], strides=[1,2,2,1], padding='SAME')
    conv2_bn = tf.layers.batch_normalization(conv2_pool)
    conv2_bn = tf.nn.dropout(conv2_bn, 0.8)

    conv3 = tf.nn.conv2d(conv2_bn, conv3_filter, strides=[1,1,1,1], padding='SAME')
    conv3_l2 = 1e-4*tf.nn.l2_loss(conv3)
    l2.append(conv3_l2)
    conv3 = tf.nn.relu(conv3)
    conv3_bn = tf.layers.batch_normalization(conv3)

    conv4 = tf.nn.conv2d(conv3_bn, conv4_filter, strides=[1,1,1,1], padding='SAME')
    conv4_l2 = 1e-4*tf.nn.l2_loss(conv4)
    l2.append(conv4_l2)
    conv4 = tf.nn.relu(conv4)
    conv4_pool = tf.nn.max_pool(conv4, ksize=[1,2,2,1], strides=[1,2,2,1], padding='SAME')
    conv4_bn = tf.layers.batch_normalization(conv4_pool)
    conv4_bn = tf.nn.dropout(conv4_bn, 0.7)
```

```

conv5 = tf.nn.conv2d(conv4_bn, conv5_filter, strides=[1,1,1,1], padding='SAME')
conv5_l2 = 1e-4*tf.nn.l2_loss(conv5)
l2.append(conv5_l2)
conv5 = tf.nn.relu(conv5)
conv5_bn = tf.layers.batch_normalization(conv5)

conv6 = tf.nn.conv2d(conv5_bn, conv6_filter, strides=[1,1,1,1], padding='SAME')
conv6_l2 = 1e-4*tf.nn.l2_loss(conv6)
l2.append(conv6_l2)
conv6 = tf.nn.relu(conv6)
conv6_pool = tf.nn.max_pool(conv6, ksize=[1,2,2,1], strides=[1,2,2,1], padding='SAME')
conv6_bn = tf.layers.batch_normalization(conv6_pool)
conv6_bn = tf.nn.dropout(conv6_bn, 0.6)

flat = tf.contrib.layers.flatten(conv6_bn)

out = tf.contrib.layers.fully_connected(inputs=flat, num_outputs=10, activation_fn=None)
return out, l2

```




Figure 3: accuracy 84%

In addition, I looked into some papers about Cifar10 training but cannot reproduce their networks due to my limited GPU memory. They have accuracy about 93%. For example: <https://arxiv.org/pdf/1412.6806.pdf>.

As for the visualization of layer output, my laptop has some issue with tensorboard and it cannot display `tf.image.summary` at this moment. I apologize for this.